Trajectory A Assessment | More

## Trajectory A Assessment

### Dimension 0 (A)

Seeded test suite execution outcome. Read-only. Flag concerns if results seem inconsistent.

#### Dimension 0 · Goal completeness (A)

Seeded hidden-test result. PASS is 100%, FAIL is below 100%. Read-only. Flag questionable results instead of editing.

PASS

### Trajectory A Dimensions

Score dimensions 1 to 9 as NO_ISSUE, MINOR_ISSUE, or MAJOR_ISSUE with citations. Score final overall quality from 1 to 5.

#### 1 · Correctness (A)*

Evaluate whether the implementation satisfies the material stated behavior, relevant edge cases, regression expectations, and non-execution requirements.

Assign **MAJOR_ISSUE** when a material requirement is incorrect or missing, a significant regression exists, or the implementation requires substantial user repair. This rating also applies when the solution relies on a fake or deceptive fix, such as hardcoding, swallowed failures, or weakened tests.

Assign **MINOR_ISSUE** when the core behavior is correct but a localized, non-blocking correctness gap or edge-case defect remains.

Assign **NO_ISSUE** when the implementation satisfies prompt requirements and no substantive correctness defect is identified.

If a model has a **FAIL** overall result, then Dimension 1 should never be labeled **NO_ISSUE**.

Select...

#### 1 · reason (A)*

#### 2 · Architecture, Modularity & Trade-off Soundness (A)*

Evaluate how the solution is structured, whether its abstractions are appropriate, and whether its algorithms and data structures suit the task. Judge the design trade-offs and expected complexity in the context of the task type and scope.

Assign **MAJOR_ISSUE** when architectural weaknesses materially affect the solution. Examples include severely broken structure, circular dependencies, excessive coupling, poor abstractions, or a naive strategy on a critical execution path. An irrational or mismatched design, including extreme over-engineering, also belongs here when it degrades performance or usability.

Assign **MINOR_ISSUE** when the design is functional but has localized weaknesses, such as being somewhat ad hoc, inefficient, over-engineered, or under-designed.

Assign **NO_ISSUE** when the design has clean boundaries, sound trade-offs, and appropriate extensibility, with no substantive architectural defect.

Select...

2 · reason (A)*

Cite specific steps, files, or diff lines.

0 / ≥ 15 words

3 · Codebase Conventions, Utility Reuse & Diff Discipline (A)*

Evaluate whether the implementation follows the project idioms, reuses existing helpers, and avoids unnecessary duplication. Consider whether the patch stays focused on the requested work and leaves the code appropriately cleaned up.

Assign **MAJOR_ISSUE** when convention violations, reinvention, or patch noise create a material weakness. This includes ignoring repository conventions, duplicating core infrastructure, or changing unrelated files without justification when the impact meets the major-issue threshold.

Assign **MINOR_ISSUE** for localized inconsistency, duplication, or unnecessary churn.

Assign **NO_ISSUE** when the implementation follows the codebase conventions, reuses its major utilities, and introduces minimal noise.

Select...

3 · reason (A)*

Cite specific steps, files, or diff lines.

0 / ≥ 15 words

4 · Robustness, Safety & API Stability (A)*

Evaluate how the implementation handles failures, malformed inputs, resource use, concurrency, security, and compatibility. Judge failure behavior in context. Strict fail-fast behavior may be appropriate for contract or data-integrity violations, while safe graceful degradation may be appropriate for user-facing resilience. Consider whether the approach increases the likelihood or severity of a given edge case when evaluating whether the agent should have specifically tested for it.

Assign **MAJOR_ISSUE** when fragile behavior, masked errors, unsafe exception handling, or unplanned breaking changes materially affect the solution. Severe bugs, deadlocks, resource leaks, vulnerabilities, silent corruption, and unhandled API breakage are examples of weaknesses covered by this rating.

Assign **MINOR_ISSUE** when primary cases are handled but localized gaps remain in rare edge conditions or compatibility.

Assign **NO_ISSUE** when the implementation provides appropriate safeguards, handles expected failures soundly, and preserves required APIs and compatibility.

Select...

4 · reason (A)*

Cite specific steps, files, or diff lines.

0 / ≥ 15 words

### 5 · Instruction Following & Scope Discipline (A)*

Evaluate whether the agent follows explicit requirements, prohibitions, scope boundaries, and workflow rules.

Assign **MAJOR_ISSUE** when the agent blatantly ignores core constraints or introduces extensive additional scope.

Assign **MINOR_ISSUE** when the core requirements are met but a minor instruction is missed or the scope is slightly stretched.

Assign **NO_ISSUE** when every material instruction is followed without scope creep.

Select...

### 5 · reason (A)*

Cite specific steps, files, or diff lines.

0 / ≥ 15 words

### 6 · Planning, Research & Tool Proficiency (A)*

Evaluate how the agent breaks down the task, navigates the codebase, explores possible solutions, uses tools, and adapts to what it learns. Research and exploration should be proportionate to the task. A localized change should allow rapid convergence, while a complex architectural task may require broader exploration of alternatives.

Assign **MAJOR_ISSUE** when aimless loops, blind edits, repeated rereading, or an inability to make progress constitute a material process weakness.

Assign **MINOR_ISSUE** when the agent reaches the correct area but takes avoidable detours or commits prematurely to an approach.

Assign **NO_ISSUE** when the agent follows a clear plan, navigates systematically, uses tools effectively, and considers the primary alternative.

Identify thrashing when the agent repeats an action three or more times without pivoting strategy. Examples include repeating the same test or command without material adjustments, or repeatedly reading the same files without extracting new information.

Select...

### 6 · reason (A)*

Cite specific steps, files, or diff lines.

0 / ≥ 15 words

### 7 · Verification & Testing Discipline (A)*

Evaluate the agent verification evidence, edge-case coverage, preservation of test integrity, debug hygiene, and honesty about results. Test weakening must always be labeled as a MAJOR_ISSUE.

Assign **MAJOR_ISSUE** when the agent performs no meaningful verification, ignores failing output, weakens or deletes tests, disables assertions, or hardcodes expected results. Weak checks, missed obvious cases, debug artifacts, or inadequately explained failures also warrant this rating when they meet the general major-issue threshold.

Assign **MINOR_ISSUE** when the provided verification is used adequately but testing remains mostly ad hoc or adds little coverage of relevant edge cases, leaving a localized weakness.

Assign **NO_ISSUE** when the agent performs appropriate functional verification and uses sound tests without weakening them.

Select...

**7 · reason (A)***

Cite specific steps, files, or diff lines.

0 / ≥ 15 words

**8 · Intermediate Communication & Risk Explanation (A)***

Evaluate communication during the trajectory, including explanations of risks, milestone updates, blocker reporting, and the frequency of communication.

Assign **MAJOR_ISSUE** when the agent remains silent during dangerous operations or conceals critical blockers until the end. Missing important risk warnings, communicating poorly over long stretches, or producing repetitive and unhelpful updates also belongs here when the weakness is material.

Assign **MINOR_ISSUE** when intermediate communication is adequate but has a localized weakness in being too brief or too verbose.

Assign **NO_ISSUE** when material risks and progress are communicated clearly, with no substantive communication defect.

Select...

**8 · reason (A)***

Cite specific steps, files, or diff lines.

0 / ≥ 15 words

**9 · Final Summary & Presentation Quality (A)***

Evaluate the final response for factual accuracy, clarity, completeness, verification reporting, and honesty about what is known.

Assign **MAJOR_ISSUE** when the summary is materially vague, incomplete, disorganized, or inconsistent with the actual work. This rating also applies when the response invents or misrepresents work or verification results, is deceptive, or provides no meaningful final summary.

Assign **MINOR_ISSUE** when the final summary is factual and usable but contains a minor omission or insufficient explanation.

Assign **NO_ISSUE** when the final summary is clear and accurate.

Select...

**9 · reason (A)***

Cite specific steps, files, or diff lines.

0 / ≥ 15 words

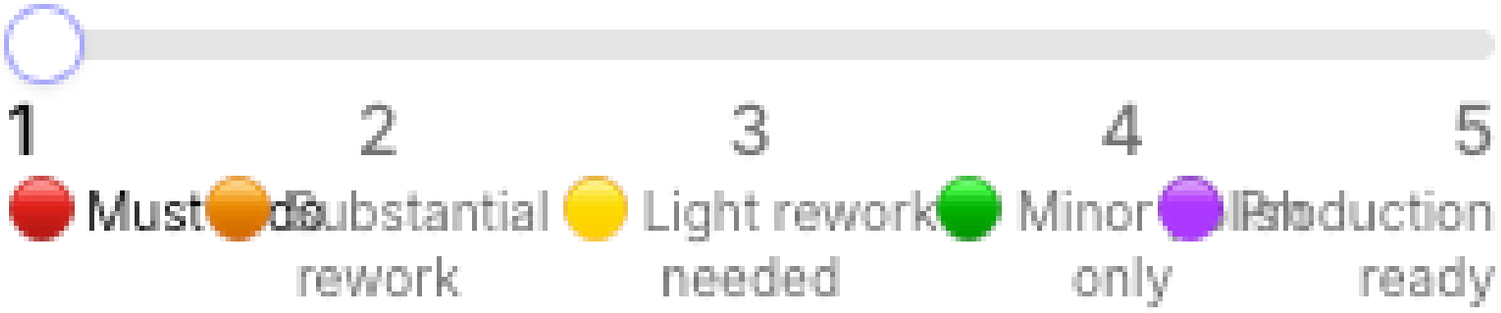
**Final quality score (A)***

Assign overall quality from 1 to 5. Explain your score and ensure it aligns with your pairwise verdict.

1 Must redo — 2 Substantial rework — 3 Light rework needed — 4 Minor only — 5 Production ready

## Final quality score (A)*

Assign overall quality from 1 to 5. Explain your score and ensure it aligns with your pairwise verdict.




## Final quality explanation (A)*

Provide evidence-based rationale for the assigned 1 to 5 quality score.

## Trajectory B Assessment

### Dimension 0 (B)

Seeded test suite execution outcome. Read-only. Flag concerns if results seem inconsistent.

#### Dimension 0 · Goal completeness (B)

Seeded hidden-test result. PASS is 100%, FAIL is below 100%. Read-only. Flag questionable results instead of editing.

FAIL

### Trajectory B Dimensions

Score dimensions 1 to 9 as NO_ISSUE, MINOR_ISSUE, or MAJOR_ISSUE with citations. Score final overall quality from 1 to 5.

#### 1 · Correctness (B)*

Evaluate whether the implementation satisfies the material stated behavior, relevant edge cases, regression expectations, and non-execution requirements.

Assign **MAJOR_ISSUE** when a material requirement is incorrect or missing, a significant regression exists, or the implementation requires substantial user repair. This rating also applies when the solution relies on a fake or deceptive fix, such as hardcoding, swallowed failures, or weakened tests.

Assign **MINOR_ISSUE** when the core behavior is correct but a localized, non-blocking correctness gap or edge-case defect remains.

Assign **NO_ISSUE** when the implementation satisfies prompt requirements and no substantive correctness defect is identified.

If a model has a **FAIL** overall result, then Dimension 1 should never be labeled **NO_ISSUE**.

Select...

#### 1 · reason (B)*

#### 2 · Architecture, Modularity & Trade-off Soundness (B)*

Evaluate how the solution is structured, whether its abstractions are appropriate, and whether its algorithms and data structures suit the task. Judge the design trade-offs and expected complexity in the context of the task type and scope.

Assign **MAJOR_ISSUE** when architectural weaknesses materially affect the solution. Examples include severely broken structure, circular dependencies, excessive coupling, poor abstractions, or a naive strategy on a critical execution path. An irrational or mismatched design, including extreme over-engineering, also belongs here when it degrades performance or usability.

Assign **MINOR_ISSUE** when the design is functional but has localized weaknesses, such as being somewhat ad hoc, inefficient, over-engineered, or under-designed.

Assign **NO_ISSUE** when the design has clean boundaries, sound trade-offs, and appropriate extensibility, with no substantive architectural defect.

Select...

### 2 · reason (B)*

Cite specific steps, files, or diff lines.

0 / ≥ 15 words

### 3 · Codebase Conventions, Utility Reuse & Diff Discipline (B)*

Evaluate whether the implementation follows the project idioms, reuses existing helpers, and avoids unnecessary duplication. Consider whether the patch stays focused on the requested work and leaves the code appropriately cleaned up.

Assign **MAJOR_ISSUE** when convention violations, reinvention, or patch noise create a material weakness. This includes ignoring repository conventions, duplicating core infrastructure, or changing unrelated files without justification when the impact meets the major-issue threshold.

Assign **MINOR_ISSUE** for localized inconsistency, duplication, or unnecessary churn.

Assign **NO_ISSUE** when the implementation follows the codebase conventions, reuses its major utilities, and introduces minimal noise.

Select...

### 3 · reason (B)*

Cite specific steps, files, or diff lines.

0 / ≥ 15 words

### 4 · Robustness, Safety & API Stability (B)*

Evaluate how the implementation handles failures, malformed inputs, resource use, concurrency, security, and compatibility. Judge failure behavior in context. Strict fail-fast behavior may be appropriate for contract or data-integrity violations, while safe graceful degradation may be appropriate for user-facing resilience. Consider whether the approach increases the likelihood or severity of a given edge case when evaluating whether the agent should have specifically tested for it.

Assign **MAJOR_ISSUE** when fragile behavior, masked errors, unsafe exception handling, or unplanned breaking changes materially affect the solution. Severe bugs, deadlocks, resource leaks, vulnerabilities, silent corruption, and unhandled API breakage are examples of weaknesses covered by this rating.

Assign **MINOR_ISSUE** when primary cases are handled but localized gaps remain in rare edge conditions or compatibility.

Assign **NO_ISSUE** when the implementation provides appropriate safeguards, handles expected failures soundly, and preserves required APIs and compatibility.

Select...

### 4 · reason (B)*

Cite specific steps, files, or diff lines.

0 / ≥ 15 words

### 5 · Instruction Following & Scope Discipline (B)*

Evaluate whether the agent follows explicit requirements, prohibitions, scope boundaries, and workflow rules.

Assign **MAJOR_ISSUE** when the agent blatantly ignores core constraints or introduces extensive additional scope.

Assign **MINOR_ISSUE** when the core requirements are met but a minor instruction is missed or the scope is slightly stretched.

Assign **NO_ISSUE** when every material instruction is followed without scope creep.

Select...

### 5 · reason (B)*

Cite specific steps, files, or diff lines.

0 / ≥ 15 words

### 6 · Planning, Research & Tool Proficiency (B)*

Evaluate how the agent breaks down the task, navigates the codebase, explores possible solutions, uses tools, and adapts to what it learns. Research and exploration should be proportionate to the task. A localized change should allow rapid convergence, while a complex architectural task may require broader exploration of alternatives.

Assign **MAJOR_ISSUE** when aimless loops, blind edits, repeated rereading, or an inability to make progress constitute a material process weakness.

Assign **MINOR_ISSUE** when the agent reaches the correct area but takes avoidable detours or commits prematurely to an approach.

Assign **NO_ISSUE** when the agent follows a clear plan, navigates systematically, uses tools effectively, and considers the primary alternative.

Identify thrashing when the agent repeats an action three or more times without pivoting strategy. Examples include repeating the same test or command without material adjustments, or repeatedly reading the same files without extracting new information.

Select...

### 6 · reason (B)*

Cite specific steps, files, or diff lines.

0 / ≥ 15 words

### 7 · Verification & Testing Discipline (B)*

Evaluate the agent verification evidence, edge-case coverage, preservation of test integrity, debug hygiene, and honesty about results. Test weakening must always be labeled as a MAJOR_ISSUE.

Assign **MAJOR_ISSUE** when the agent performs no meaningful verification, ignores failing output, weakens or deletes tests, disables assertions, or hardcodes expected results. Weak checks, missed obvious cases, debug artifacts, or inadequately explained failures also warrant this rating when they meet the general major-issue threshold.

Assign **MINOR_ISSUE** when the provided verification is used adequately but testing remains mostly ad hoc or adds little coverage of relevant edge cases, leaving a localized weakness.

Assign **NO_ISSUE** when the agent performs appropriate functional verification and uses sound tests without weakening them.

Select...

### 7 · reason (B)*

Cite specific steps, files, or diff lines.

0 / ≥ 15 words

### 8 · Intermediate Communication & Risk Explanation (B)*

Evaluate communication during the trajectory, including explanations of risks, milestone updates, blocker reporting, and the frequency of communication.

Assign **MAJOR_ISSUE** when the agent remains silent during dangerous operations or conceals critical blockers until the end. Missing important risk warnings, communicating poorly over long stretches, or producing repetitive and unhelpful updates also belongs here when the weakness is material.

Assign **MINOR_ISSUE** when intermediate communication is adequate but has a localized weakness in being too brief or too verbose.

Assign **NO_ISSUE** when material risks and progress are communicated clearly, with no substantive communication defect.

Select...

### 8 · reason (B)*

Cite specific steps, files, or diff lines.

0 / ≥ 15 words

### 9 · Final Summary & Presentation Quality (B)*

Evaluate the final response for factual accuracy, clarity, completeness, verification reporting, and honesty about what is known.

Assign **MAJOR_ISSUE** when the summary is materially vague, incomplete, disorganized, or inconsistent with the actual work. This rating also applies when the response invents or misrepresents work or verification results, is deceptive, or provides no meaningful final summary.

Assign **MINOR_ISSUE** when the final summary is factual and usable but contains a minor omission or insufficient explanation.

Assign **NO_ISSUE** when the final summary is clear and accurate.

Select...

### 9 · reason (B)*

Cite specific steps, files, or diff lines.

0 / ≥ 15 words

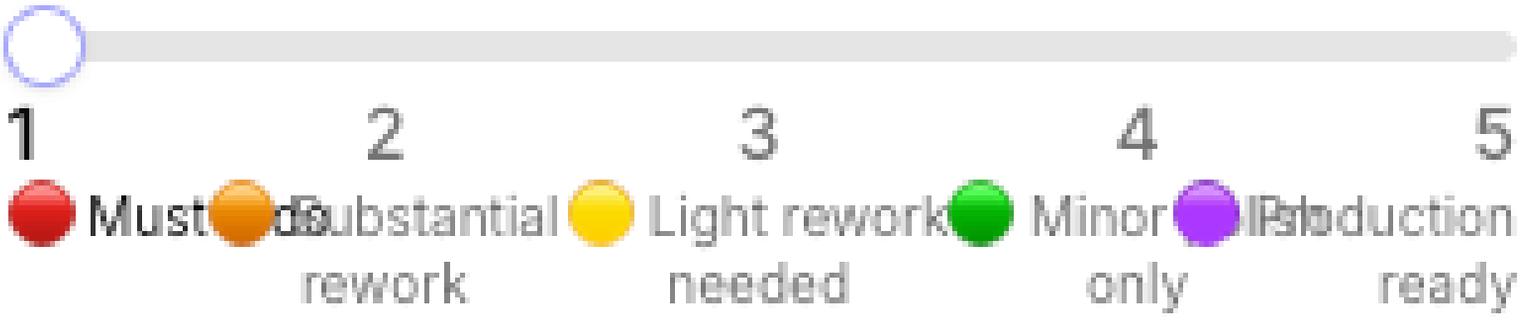
### Final quality score (B)*

Assign overall quality from 1 to 5. Explain your score and ensure it aligns with your pairwise verdict.

1 Must do
2 Substantial rework
3 Light rework needed
4 Minor only
5 Production ready

## Final quality score (B)*

Assign overall quality from 1 to 5. Explain your score and ensure it aligns with your pairwise verdict.




## Final quality explanation (B)*

Provide evidence-based rationale for the assigned 1 to 5 quality score.

## Pairwise Evaluation

### Pairwise Evaluation*

Select your pairwise preference rating between Trajectory A and Trajectory B.

Select...

A MUCH BETTER

A BETTER

TIE

B BETTER

B MUCH BETTER

## Pairwise Evaluation

### Pairwise Evaluation*

Select your pairwise preference rating between Trajectory A and Trajectory B.

Select...

### Comparative Rationale*

Explain your pairwise preference rating.